

# CS140E: embedded OS

Stanford Winter 2026

# The Staff

- Dawson Engler (if stuff is broken it's my fault)
- Head TA: Joseph Shetaye (rockets, booms, winning bets)
- TAs:
  - Maximilien Cura (Rust + insane hacks)
  - Aditya Sriram (ox64 riscv + bass)
- Section leads: (All: unpaid volunteers(!))
  - Rohan Ram Chanani + Asanshay Gupta (the GPU guys)
  - James Yu-tang Chen + Sai Ketan Konkimalla (the DOOM! guys)
  - Ron Dubinsky (Paxos+networking)
  - Stuart Sul (ELF, real world, and top-tier AI)
  - (Almost 10% of 140e'25 came back to help!)

# The Students

Some staff-centric context:

- Lab classes usually capped around 30 students: We graduated 73 last year.
- Class start: 530pm. Common: staff helping til after midnight.
- Our former head TA Joe Tan sighted many times at 3am.

This year:

- 180 applications.
- Very hard folders: already a very self-selecting group.
- Max, Joseph, Asanshay, Ron each spent roughly 20-30hrs over 6 weeks to pick ~80 folders.
- (Context: More hours than a typical ms/phd admission committee member.)


## A note about work

Stanford doesn't pay to staff lab classes:

- Most of the staff are unpaid volunteers.
- This class is way too much work even if paid.
- So if you need help outside class please only bug non-volunteers (Max, Aditya, me).

# Class: Write small, clean OS on a r/pi zero w.

[https://en.wikipedia.org/wiki/Raspberry\\_Pi](https://en.wikipedia.org/wiki/Raspberry_Pi)



# Class philosophy

Write a complete, narrow OS all the way down to the bare metal.

Typical OS class:

- a bunch of complicated stuff covered superficially, but with a lot of bugs and starter code.

Our approach:

- You'll write complete, simple versions that work. On real hardware.
- In fact: b/c of novel tools you'll build: it will be *surprising* if code broken.
- And: Because you write it all, you will understand much more thoroughly
- The hope: easy to do delta off of your knowledge to more fancy things

# Why write bare-metal OS code?

## *Power:*

- Real control of real hardware = real superpower.
- Can build many things easily that are essentially impossible on linux/macos.
- Guaranteed nanosecond performance, actually secure, many devices.

## *Powerful*

- If you can write kernel code: can write almost anything (non-math-y)
- Bugs here are some of the hardest yet invented by entropy:
  - device hardware errata, page table mistakes, interrupt lock-ups, trashing a register during context switching = good luck with that.
- You'll think and code more clearly b/c otherwise it sucks.
- Bonus: the bugs you hit will *really* teach you about the

# Why I do it.

Small bare-metal system you wrote yourself:

- Fluency is 100.
- Flow state is maxxed.
- Very very easy to drop in and immediately make changes.
- And since it's bare metal, you can make it do crazy things, quickly.
- At my age, not as easy to have fun writing code. This does it :)

# Class org: Labs. More labs. Then: Final project.

Organization:

- Very little lecture (today is longest)
- 18 labs, 2x per week. Start at 530pm. Go til you want to leave.
- Prelab material (plus gradescope prelab) before Lab.
- Lab: walk in, by the end of the lab, you have a complete working simple version  
of a key trick.
- E.g., 300-400 line working interrupt handling or veritual memory system.

# This class has no textbook.

Readings = raw primary sources:

- datasheets and architecture manuals.
- (Though we do give you annotations and commentary.)
- Some are intense: virtual memory is about 100 pages of tough material.

# This class has no textbook.

Why:

- Classes present a fake reality. The real real world is not:
  - A clean textbook of systematized knowledge.
  - With pretty, worked out problems
  - And bolded text saying you should do multigrid relaxation on problem 4.
- It is:
  - Difficult to understand documents with no high bit. That are: wrong, incomplete, not written to be used ("passive definitional voice"), describing nouns that themselves have bugs.

# This class has no textbook.

Promises:

1. If you understand these, you can do things others cannot.
2. After 10 datasheets, the 11th is not a big deal.

# Who is this class for?

Class is heavily tuned for those that are interested in low level code/hw.

- It is built to move fast with a bunch of people that are very interested
- Only take the class if:
  - i. You find Stanford easy and/or
  - ii. You are very interested and have a very open quarter.
- Sub-category: If you're on the spectrum (hi).

Our goals:

- We don't waste your time;
- You cover adult stuff in a way and level you haven't seen;
- You'll find many hardcore people you want to work with in the future

**Warning:** Don't take this class thinking it is an "easier version of 112" (it is not).

# Why R/Pi A+?

- Most OSes write code on a fake simulator
  - A lot of work, not that cool at the end
- R/pi = real computer for about \$16 and an ounce of weight
- Many examples / blog posts of how to do various things
- Unlike most machines, makes interacting with the real world easy
- Can build many interesting systems b/c can use weird hardware easily
  - motion sensors, IR sensors, accelerometer, gyroscope, light sensor...
- Actually fast compared to modern SBCs (1GHz), and has an MMU.

# **Goal: you will develop two super-powers**

## **Power 1: Differential debugging**

- Efficiently answering "why doesn't this work" for complex things
- Swap working pieces + Binary search

## **Power 2: Epsilon development**

- Epsilon sprint paradox: When building systems, the smaller the step you take, the faster you can run.

# Differential debugging

You write code, it doesn't work, the error could be:

- The code you wrote
- Hardware fault (bad manufacturing, smoked something)
- Wiring mistake
- Subtle cache issue
- Compiler problem (more on this)
- ...

You will get good at breaking down problems by swapping pieces between a **working system** (yesterday's code, your partner's lab) and a **non-working system** (today's code, your lab)

# Example from next lab

You get the following set of stuff:

To run you:

1. Copy blink.bin to sd
2. Wire up led
3. Wire up serial device
4. Plug into your laptop




**It doesn't work. But your partner's does.**

# What is messed up?

Partner's	Yours
Working	Not working


What to do first?



# What is messed up?

Partner's	Yours
Working	Not working

- What does swapping tell you if it doesn't work?
- What does swapping tell you if it works?



# What is messed up?

Partner's	Yours
Working	Not working

Swapping works: how to narrow down with least work?




# What is messed up?

Partner's	Yours
Working	Not working

**Entire class:** whenever we control device, has some software component S (can be wrong) and some hardware component H (can be broken).

Doesn't work = linear equation solving with two variables. **How to isolate?**



# Epsilon sprinting: Slow is fast

What is wrong?

- If I did X, it's X.
- If I did  $X_1 + X_2 + \dots + X_n$  it could be any, or some combination.

Inverting crash/bug to root cause is **much harder** in the latter case.

**My epsilon-sprint theorem:**

Given a working system  $W_k$  and a change  $C$ , then as  $|C| \rightarrow \varepsilon$ , the time + computation (IQ) it takes to figure out why  $\{ W_k + C \}$  doesn't work goes to 0.

**Related claim:** the time it takes to debug why a change broke the system increases non-linearly with the size of the change.

# Before TA: make sure you did an epsilon delta

Assume: 3 10-minute bugs per 80 students and 6 TAs:

- Arrival rate:  $(N \times \text{bugs\_per\_student}) / \text{hours}$ :  $(80 \times 3) / 6 = 40 \text{ bugs/hour}$
- Service rate: 6 bugs per TA / hr.
- Service capacity:  $6 \text{ TAs} * 6 \text{ bugs / hr} = 36/\text{hr}$
- Utilization:  $\text{Arrival} / (\text{C} \times \text{U}) = 40/36 = 110\%$ .
  - After 6 hours: 24 deep. (Worse b/c TAs quit, bursty)

140e code development:

- PLEASE DO NOT DO multiple changes and then ask "it doesn't work"
- Always take working system, add the smallest epsilon change possible to make break. Only ask then.
- Useful IRL. Also, often: you fix bug.

# Administrivia

## Rules:

- Feel free to leave early, but don't be late (makes us run O(n))
- Don't use llms or other people's code. You have to write.
- Don't miss more than 2 labs (even that is very tough).

## Grade:

- "A" requires 3+ hard extensions
- Absolutely must turn in lab w/in 7 days. No exceptions!
- Usually prelab due before lab. Won't accept after.
- Will have a final project (roughly 3+ labs worth of work)

**Plus side:** We pay for food. We pay for equipment. We stay til midnight or so.

# Administrivia

Two labs each week.

- Each lab will have pre-lab work you should turn in before lab
- Ideally finish during the lab period (I will stay til everyone is done)
- Must finish within a week of the lab, or start losing a letter grade each day
- Must pre-arrange missed labs. It's a problem to miss more than a couple.

There (tentatively) will be three "capstone" homework assignments that consolidate a chunk of labs together.

If you've done the lab, this shouldn't be a big deal.

# New: grade partly depends on understanding

## *Problem:*

- If we don't look for cheating, it can happen.
- Because we've been able to pick great people, it is rare. But not 0.

## *DO not want:*

- To ignore: unfair.
- But also don't want: in-class exams, cold-calling, MOSS detection.

## *Hack*

- Grade partly based on how well you understand your code *when done*.
- It's fine to be confused before. But after should have a grasp of it.
- Pro: no tests, more time for helping people.

# Administrivia

You can work with other people!

However, you *must* type and turn in everything yourself.

(We will pay for food you order during lab.)

# What to do now

1. Clone the class git repository:

```
git clone git@github.com:dddrreee/cs140e-26win.git
```

<https://github.com/dddrreee/cs140e-26win>

2. For lab next Tuesday (one week), make sure you:

- Have a way to write either a micro-SD or SD card
- Have a way to plug in a standard USB device
- Do PRELAB for lab 2-trust



